

Report On

SmartSeg: Customer Segmentation using K-means

Submitted in partial fulfillment of the requirements of the Mini project in
Semester IV of Second Year Computer Science Engineering (Data Science)

by
Rohit Khanolkar (Roll No. 11)
Taha Khan (Roll No. 10)
Parvez Khan (Roll No. 09)
Meet Kevadiya (Roll No. 08)

Mentor
Prof. Bhavika Joshi



University of Mumbai

Vidyavardhini's College of Engineering & Technology

Department of Computer Science Engineering (Data Science)



(A.Y. 2025-26)

Abstract

Customer segmentation is an important process in today's business analytics. It helps organizations understand customer behavior and make better decisions. This project presents a web-based customer segmentation system that uses the K-Means clustering algorithm to group customers based on key traits like age, annual income, and spending habits. The system starts with data preprocessing, which includes cleaning and normalizing the data to ensure consistency and improve clustering quality. The optimal number of clusters is determined using the Elbow Method. This method balances accuracy and efficiency.

The trained model is integrated into a user-friendly web application. Users can input customer details and get real-time segmentation results. The system classifies customers into meaningful groups, such as high-value customers, low-spending customers, and inactive users. This enables businesses to create targeted marketing strategies and enhance customer engagement. Using machine learning improves the speed, accuracy, and scalability of the segmentation process compared to traditional manual methods.

Additionally, the project shows how data mining and machine learning techniques can solve real-world problems. The results indicate that the system can reliably identify patterns within customer data and generate dependable clusters. Overall, this project emphasizes the effectiveness of combining machine learning with web technologies to provide an accessible, efficient, and scalable solution for customer segmentation.

Vidyavardhini's College of Engineering & Technology
Department of Computer Science Engineering (Data Science)

CERTIFICATE

This is to certify that the Mini Project entitled “ **SmartSeg: Customer Segmentation using K-means**” is a bonafide work of **Rohit Khanolkar (Roll No. 11), Taha Khan (Roll No. 10), Parvez Khan (Roll No.09), Meet Kevadiya (Roll No. 08)** submitted to the University of Mumbai in partial fulfillment of the requirement for the award of the degree of “**Bachelor of Engineering**” in Semester IV of Second Year “**Computer Science Engineering (Data Science)**”.

Prof. Bhavika Joshi
Mentor

Prof. Satish Salunkhe
(Ph.D.)
Head of Department

Dr. Rakesh Himte
Principal

Vidyavardhini's College of Engineering & Technology
Department of Computer Science Engineering (Data
Science)

Mini Project Approval

This Mini Project entitled “ **SmartSeg: Customer Segmentation using K-means**” by **Rohit Khanolkar (Roll No. 11), Taha Khan (Roll No. 10), Parvez Khan (Roll No.09), Meet Kevadiya (Roll No. 08)** is approved for the degree of **Bachelor of Engineering** in Semester IV of Second Year **Computer Science Engineering (Data Science)**.

Examiners

1.....
(Internal Examiner Name & Sign)

2.....
...
(External Examiner Name & Sign)

Date:

Place:

Contents

Abstract	i
Acknowledgments	ii
List of Abbreviations	iii
List of Figures	iv
List of Tables	v
List of Symbols	vi
1 Introduction	1
1.1 Introduction	1
1.2 Problem Statement & Objectives	3
1.3 Scope	4
2 Literature Survey	6
2.1 Survey of Existing System/SRS	6
2.2 Limitation Existing system or Research gap	7
2.3 Mini Project Contribution	10
3 Proposed System (eg New Approach or Data Summarization)	12
3.1 Introduction	12
3.2 Architecture/ Framework/Block diagram	13
3.3 Algorithm and Process Design	15
3.4 Details of Hardware & Software	18
3.5 Experiment and Results for Validation and Verification	19
3.6 Analysis	24
3.7 Conclusion and Future work.	25
References	27
4 Annexure	
4.1 Published Paper /Camera Ready Paper/ Business pitch/proof of concept	

Chapter 1

Introduction

1.1. Problem Statement:

Despite having access to massive volumes of customer details, numerous companies find it difficult to turn that information into useful insights. Generated daily from e-commerce activities, physical stores, and web-based services, the flow of data continues growing. It captures elements like age, buying habits, choices, and how much people spend. Because of disconnected tools or unclear strategies, potential often remains untapped. What seems abundant on the surface frequently lacks clarity underneath.

Still, spotting real trends in customer information remains tough for many groups. When insights get missed, grasping what people want, like, or do grows harder. Because of this gap, businesses tend to fall back on broad campaigns that miss their mark more than hit it. Sooner or later, mismatched outreach chips away at trust - satisfaction dips, results follow.

Most older ways to sort customers rely on fixed rules or hands-on work, making them slow when dealing with massive amounts of information. Handling such volume becomes harder because these approaches demand too much time while increasing chances for mistakes. Flexibility suffers as customer habits shift over time, leaving outdated models behind. Hidden trends within data remain unnoticed since those techniques cannot detect subtle groupings automatically.

Machine learning opens doors to automated data analysis, uncovering hidden trends without constant human oversight. Still, smaller companies often struggle - tools built for such tasks remain out of reach, too complex or poorly designed for everyday use. One reason behind building smart tools lies in sorting people by traits like age or earnings without manual effort. Efficiency matters - so does growth potential and simplicity when putting these systems to work. Decisions backed by information

become possible, shifting how promotions are planned over time.

This work tackles the issue through a web tool built on machine learning, designed to sort customers clearly and efficiently. While focusing on usability, it applies smart algorithms behind the scenes without demanding technical skill. Instead of complex setups, the system runs in browsers, making access straightforward. Because automation helps reduce manual effort, results appear faster than traditional methods. Even so, accuracy remains central to how segments are formed. Through real-time processing, patterns emerge more naturally across diverse user groups.

1.2. Objectives

Project goals appear next. What comes now is a list of aims. Listed here are the targets.

You see before you the intentions. These represent what needs accomplishing.

Shown below are purposes meant to guide work.

- To study the concept of customer segmentation and its importance in business analytics
- To understand and implement machine learning techniques for data analysis
- To apply K-Means clustering algorithm for grouping customers
- Examining consumer information through characteristics like age, household earnings, then spending habits shapes clearer understanding. Patterns emerge when viewing these factors together yet independently. A person's financial behavior often shifts with life stage.

Income level might influence purchase choices more than expected.

Spending patterns reveal tendencies beyond what age alone suggests. Insights grow sharper focusing on combined variables instead of isolated ones

- To design and develop a web-based interface for user interaction
- To integrate frontend and backend technologies with the machine learning model
- Accurate forecasting of customer groups begins with clear data patterns. Efficiency emerges when models adapt to real usage trends. Prediction improves through consistent feedback loops. Segment clarity often follows simplified assumptions. Results depend heavily on initial setup quality
- For simplicity, speed matters most when designing tools meant for people without technical skills

1.3. Project Scope :

This work centers on building a tool driven by machine learning to sort customers through clustering methods. Built to examine user information, it forms groups where individuals share clear patterns. With just a few inputs - like age, yearly earnings, and how much they spend - it reveals habits in an accessible manner. Efficiency emerges not from complexity, instead from targeted analysis grounded in key traits.

Although the platform runs on a web interface, it captures client information and delivers instant grouping outcomes. Running behind the scenes, the server takes submitted data and uses K-Means Clustering to assign a category. Through an intuitive display, findings appear clearly - helping even non-technical individuals navigate the tool. With little training needed, access remains open to a broader audience despite underlying complexity.

This setup works well for modest or mid-sized operations needing fast, straightforward insights into client behavior. Because it sorts users into groups - like top contributors, typical buyers, or those showing little interest - it supports sharper planning. While not built for complexity, its clarity allows teams to adjust outreach with more precision. Decisions gain direction when patterns in spending and interaction become visible through this method.

Still, today's version focuses only on static divisions or user-provided inputs, missing live data updates. Instead of adapting on the fly, it runs using set information paired with a constant cluster count - falling short when conditions shift rapidly.

The project can be further expanded in the future by:

Using data collected instantly often means handling vast amounts at once

Location details might shape the profile further. Preferences come into play alongside past purchases. One factor builds on another, quietly refining the understanding. Purchase records add depth when combined with regional data. Insights grow sharper through layered information. Regional habits influence patterns seen across transactions. Past behavior links with personal choices in subtle ways

Using advanced machine learning algorithms for better accuracy

Developing a more interactive dashboard for data visualization

Integrating the system with business intelligence tools

This work builds a starting point for sorting customers, showing real-world business uses through machine learning methods. While it focuses on core ideas, its value lies in turning algorithms into actionable insights across departments.




Fig.1: Customer Segmentation System Flowchart

Chapter 2

Literature

Review

2.1 Survey of Existing System/SRS:

2.1.1 Review of Current Setup

Most older companies sort buyers by hand or follow fixed rules. Instead of complex tools, they pick obvious traits - like how old someone is, where they live, or how often they buy. These ways work at small scale, yet struggle when data grows too big. Hidden habits in buying choices usually go unnoticed. Efficiency drops once the volume increases beyond a certain point.

Some companies rely on spreadsheets or simple database searches to study customer information. Though common, these approaches demand a lot of time and often lead to mistakes. Human-driven sorting struggles to keep pace with shifting patterns in behavior. Precision drops sharply once data grows more intricate.

As data analytics progresses, certain platforms now apply statistics along with simple grouping approaches. Still, those tools typically demand specialized knowledge, making them difficult for people without a tech background. Besides that, instant or fast categorization outcomes might be missing entirely.

These days, tools powered by Machine Learning handle image segmentation automatically. One common method involves grouping data through algorithms like K-Means. Still, most of these solutions run separately, without connection to online systems. As a result, companies find it harder to apply them regularly in real workflows.

Limits of Current Systems

- Manual and time-consuming process
- Large datasets cannot be handled effectively
- Lack of accuracy in identifying customer patterns

- Requires technical knowledge
- Limited scalability and flexibility
- No real-time or quick prediction capability




Fig.2: Traditional vs Modern Customer Segmentation

2.1.2 Software Requirement Specification (SRS)

The Software Requirement Specification (SRS) describes the functional and non-functional requirements of the system. It acts as a blueprint for the development of the project.

A. Functional Requirements

These are the core features of the system:

- The system should allow users to input customer data such as age, income, and spending score
- The system should process and analyze the input data

- The system should apply the clustering algorithm to segment customers
- The system should display the predicted customer segment
- The system should provide a user-friendly interface
- The system should handle multiple inputs efficiently

B. Non-Functional Requirements

These define system performance and quality:

- Performance: The system should provide quick results
- Usability: The interface should be simple and easy to use
- Reliability: The system should provide consistent outputs
- Scalability: The system should handle larger datasets in future
- Security: Basic input validation should be ensured

C. Hardware Requirements

- Computer/Laptop
- Minimum 4GB RAM
- Internet connection

D. Software Requirements

- Operating System: Windows / Linux / macOS
- Programming Language: Python
- Frontend: HTML, CSS, JavaScript
- Backend: Flask / PHP
- Libraries: NumPy, Pandas, Scikit-learn

2.2 Limitation: Existing system or Research gap:

Even though many ways to segment customers exist, current tools often fall short when used outside controlled settings. Instead of adapting well to messy data, older strategies depend heavily on hand-crafted rules or visual inspection. Because they

demand constant oversight, these processes drain resources while slowing down decision-making. Efficiency drops further as volumes grow, leaving companies struggling to keep pace.

Hidden patterns in customer data frequently escape detection by current systems. Because manual methods rely on fixed criteria, shifts in how customers act go unnoticed. These rigid approaches lead to flawed groupings. What emerges rarely matches real customer clusters.

Even when today's tools include machine learning, complexity tends to block smooth operation without specialized knowledge. Interfaces frequently ignore ease of use, leaving people unfamiliar with tech struggling to engage fully. On top of that, isolation from web platforms weakens their real-world application in companies.

One key drawback involves missing real-time handling. While many current systems rely on fixed data collections, they fail to adjust when information changes over time. Because of this, responding to live shifts in user actions becomes difficult.

Most current setups struggle to mix ease of use with exact results. Though sophisticated designs can boost correctness, they tend to make things harder to manage. Yet lightweight versions miss fine detail, producing shallow outcomes instead.

Research Gap:

Clearly, existing shortcomings point toward needing something better. What's missing? A way forward that pulls together what currently stays apart

- Simplicity and ease of use
- Efficient data processing
- Accurate customer segmentation
- Integration with web-based platforms
- Accessibility for non-technical users

This work fills a missing piece through a browser tool built around K-Means Clustering for sorting customers into groups. With an emphasis on ease of use, the

design still delivers solid performance in grouping accuracy.




Fig.2: Traditional vs Mordern Customer Segmentation

2.3 Mini Project Contribution:

A small-scale effort brings value through a straightforward approach to sorting customers with machine learning tools. Rather than depending on old ways - like hand-driven reviews or complicated setups - this work builds toward a smooth online platform fit for people without coding skills. Efficiency shapes its core; accessibility guides its design.

One key outcome here: applying K-Means Clustering to sort customers by traits like age, earnings each year, and how much they spend. Through examining these details, patterns emerge - distinct segments form naturally. Because of this separation, companies gain clarity about who buys what. As behavior becomes visible, tailored strategies follow more

easily.

One major addition involves linking machine learning to web tools. Instead of separate parts, the build uses HTML, CSS, and JavaScript up front, while relying on Python (Flask) behind the scenes. Through this setup, the tool runs smoothly - yet remains easy to navigate. What results is a working system that feels straightforward despite its underlying complexity.

Despite its complexity behind the scenes, simplicity guides how people interact with it. Through a standard browser window, anyone enters client details without special training. Once submitted, information moves rapidly through backend logic before showing which group the algorithm assigns. Manual sorting becomes unnecessary thanks to automated handling. Faster results mean choices happen sooner than before.

One example shows machine learning tackling actual business challenges without high costs. Through this approach, improvement becomes possible over time - more information, stronger methods, live updates shaping what comes next. A simple structure supports such growth, leaving room for later upgrades without demanding full redesigns upfront.

Project Key Contributions:

- Development of a machine learning-based customer segmentation system
- Implementation of K-Means clustering for grouping customers
- Integration of frontend and backend technologies
- Creation of a user-friendly web interface
- Fast and efficient prediction of customer segments
- Demonstration of real-world application of machine learning

Chapter 3

Proposed System

3.1 Introduction :

A new approach builds a machine learning tool aimed at sorting customers more effectively while staying straightforward to operate. Instead of relying on older techniques, this design handles data review automatically, shaping groups that make sense. Efficiency comes through smart automation, reducing manual effort without losing accuracy.

One way firms learn about buyers involves sorting them into groups. Instead of guessing, a method built on information examines features like how old someone is, yearly pay, and what they tend to spend. Patterns emerge when these details are studied closely. Because of that, companies can adjust plans using real clues. Decisions get sharper once habits show up in numbers.

At its heart lies K-Means Clustering - a common method in Machine Learning for sorting data. Instead of random placement, it forms customer groups by reducing internal spread while pushing separate groups apart. Because of this separation principle, individuals showing alike traits land in one set. What emerges is a structure where likeness defines membership.

A web-based setup allows the tool to run directly in a browser, removing installation needs. Through an intuitive layout, individuals submit client details without complexity. Once information flows in, a pre-built algorithm interprets patterns within moments. Results appear instantly - each showing a likely group assignment based on behavior traits.

What makes the system stand out is how straightforward yet effective it is. Rather than relying on intricate analysis, it delivers fast outcomes even when used by someone without advanced expertise. Still, precision does not get lost - grouping methods help spot trends reliably within datasets.

A starting point for later upgrades, the suggested setup allows room to grow. With time, it might handle live information streams alongside bigger collections of records. Improvements could involve smarter computational methods that sharpen results while boosting speed.

This approach shows the way machine learning fits together with web tools to tackle actual challenges in commerce. Though built on digital foundations, its value emerges through clear problem-solving tasks. Where data meets interface, a functional response takes shape. Not simply theoretical, it operates within everyday business conditions. Through adaptive methods, results form steadily over time. Even without complex design, usefulness becomes visible in practice.

3.2. Architecture –




Fig.3: System Architecture of Customer Segmentation System

The layout of the suggested customer grouping tool shows its main parts and how they

interact. Starting at entry points, information moves step by step toward results. Each stage connects logically, guiding raw inputs into structured outcomes. Flow direction becomes clear when tracing steps from beginning to end.

Starting from the customer data, key details like age, yearly earnings, and spending habits form the core. Built on these features, the information moves into the learning phase of the algorithm.

After gathering information, cleaning begins - errors get fixed, gaps filled through logical estimates. Inconsistencies fade when outliers are adjusted or dropped quietly. Scaling reshapes numbers so none dominate by size alone. Normalization smooths differences across features, letting patterns emerge more clearly. Accuracy climbs when inputs arrive in balanced form. Without these steps, models often misread signals buried beneath noise.

Following cleanup, data moves into K-Means Clustering - a method sorting individuals by likeness. Though silent at first glance, it detects hidden structures across records. One by one, profiles find their place within defined segments. Each grouping emerges from subtle but consistent trends buried in behavior. Where someone lands depends entirely on how closely they mirror others nearby.

After finishing cluster analysis, save the fitted algorithm into a .pkl extension. Inside sits the acquired patterns - ready later to forecast outcomes, skipping fresh training rounds.

Following processing, a prediction component activates, linking to the trained model for output creation. Connected to this is a Flask-built backend, managing incoming user queries while enabling dialogue among interface layers and the algorithmic core.

From the web form, users enter details like age, income, or spending habits. Once submitted, that information moves to the server side. There, a saved machine learning model activates upon receiving the new values. Prediction occurs instantly based on prior training patterns. The outcome assigns each person to one specific group. Assignment depends entirely on how closely inputs match known behavior clusters.

After processing, the result shows which customer category was identified. Businesses then use this information to interpret actions and shape choices based on evidence.

3.3 Algorithm and Process Design

3.3.1 Algorithm Used

A different approach begins here - clustering customers through K-Means. This method sorts data by likeness, placing items into set groups. Though common, it relies on defining cluster count ahead of time. Grouping happens iteratively, adjusting until patterns stabilize.

Starting with a preset count of groups, the method sorts every data instance into the closest center. As the process moves forward, centers shift based on current group members. With repetition, these central points adjust repeatedly. Stability arrives when reassignment stops changing outcomes. Eventually, the structure settles without further shifts.

- K Means Algorithm Operation
 1. First comes initialization - randomly placing centroids among the data points. Then every point gets assigned to its nearest center. After that, each centroid shifts to the average position of its group. This cycle repeats until movement stops or limits trigger. Finally, clusters settle into place when changes become negligible
 2. Select the number of clusters (K)
 3. Initialize K centroids randomly from the dataset
 4. Starting with proximity, every data instance links to the closest central point. Distance typically means straight-line measurement here. Whichever center is most accessible claims the point. This process repeats per observation without exception. Closest defines assignment - nothing else influences it
 5. To find updated centroids, compute the average position of every point within a given cluster. Each group's center shifts based on its members' collective location. Using arithmetic means ensures balance across dimensions. The process repeats once per iteration cycle. Movement stops when positions stabilize between steps

6. Keep applying steps 3 and 4 repeatedly, unless the movement of the centroids becomes negligible. Once shifts fall below a set threshold, stop iterating. Progress halts when position updates are minimal. This signals stabilization has occurred across iterations. The process ends naturally at that point
7. Clusters take shape once centroids stop shifting significantly
8. Grouping individuals by shared traits happens through this method, creating clusters that make sense. These collections form when people show common patterns, shaping distinct categories naturally.

3.3.2 Process Design

- Starting with raw information, the system's structure outlines each phase needed for sorting customers. This path shows how data shifts into useful results through organized stages.
- Starting with raw entries, it moves through stages shaped by preparation steps before reaching computational analysis, ending in output formation. Each phase connects to the next through defined transitions rather than abrupt shifts, building flow without relying on repetition. Input transforms gradually under methodical adjustments until outcomes take shape naturally. Processing happens in sequence, yet each segment depends on what came earlier in quiet ways. Execution unfolds as planned, guided by prior design choices embedded long before activation.

- Process Design Steps:

Dataset → Preprocessing → K-Means → Model → User Input → Prediction → Output

- Data Collection

A set of customer data gets gathered - this includes details like how old people are, their earnings, then also a measure of how much they spend. Each piece helps shape understanding without needing everything at once.

- Data Preprocessing

After removing errors, the data gets adjusted to enhance how well the model works. Scaling follows cleaning so results become more reliable across runs.

- Model Training

A first step involves feeding cleaned data into the K-Means process for pattern grouping. Model learning happens once structured inputs are ready. Training unfolds without labels, guided by distance patterns among points. Setup relies on numeric consistency across features. Iterative refinement adjusts centroids until stability appears.

- Model Saving

A saved version of the trained model sits in a .pkl file, ready when needed. While training finishes, the output gets stored automatically. Later access happens through that same file format. Retrieval takes place whenever reuse becomes necessary. The process locks progress at that point. Nothing changes once written. A snapshot stays preserved exactly there.

- User Input

Through the web interface, a person inputs information about clients. Customer data gets added via an online form each time someone types it in.

- Prediction

A prediction begins once the system pulls in the trained model. Following that, an input gets assigned to a cluster through inference. After loading, computation determines group membership. Input arrives - then classification follows based on prior learning. The model, ready ahead of time, responds when new data enters. With everything set, grouping happens automatically upon submission.

- Output Display

A portion of customers, forecasted through analysis, appears before the viewer. The system shows who fits the expected group based on data patterns.

3.4 Details of Hardware & Software

This part outlines what tools and programs are needed to build and run the suggested customer grouping system. Such conditions help keep performance steady during operation.

3.4.1 Hardware Requirements

A single modern computer handles the workload without issue, since efficiency shapes much of the design. Though powerful machines work too, they're unnecessary - the software thrives even on basic setups.

- Processor: Intel i3 or higher
- For basic operation, a system needs at least 4 GB of memory. Better responsiveness comes with 8 gigabytes. Some tasks run more smoothly when extra capacity is available
- For storage, ensure there is no less than 20 gigabytes of unused space on the drive
- System Type: 64-bit system
- Input Devices: Keyboard and Mouse
- Output Devices: Monitor
- Average computers handle the system without trouble - most current laptops or desktops run it smoothly. Performance demands stay low, so costly gear becomes unnecessary.

3.4.2 Software Requirements

Below appears a list of software parts involved in building the system

- Operating System
Operating systems include Windows, yet also cover Linux alongside macOS
- Programming Language
Python
- Frontend Technologies

- i. HTML
 - ii. CSS
 - iii. JavaScript
- Backend Framework
 - Flask (Python-based web framework)
- Machine Learning Libraries
 - i. NumPy
 - ii. Pandas
 - iii. Scikit-learn (used for implementing K-Means Clustering)
 - iv. Database (Optional)
- Development Tools
 - i. VS Code / PyCharm
 - ii. Jupyter Notebook (for model training and testing)
 - iii. Brave Browser

3.5 Experiment and Results for Validation and Verification

This part describes tests carried out to assess how well the suggested customer grouping method works. Its purpose: confirming the system sorts clients accurately according to provided information, while delivering useful outcomes.

3.5.1 Experiment Setup

Despite its simplicity, the framework underwent evaluation with real user data featuring details like age, yearly earnings, and expenditure patterns. Following preparation steps, the information fed into a grouping mechanism built around the K-Means method. While some noise existed initially, refinement improved input quality before training began.

Following training, the system stored the model then connected it to the web platform for live forecasts. The process began by choosing K - the ideal cluster count - through the

elbow technique.




Fig.4: Elbow Method For Optimal Clusters

3.5.2 Experimental Procedure

The following steps were carried out during the experiment:

- Step 1: Load the dataset into the system
- Step 2: Perform data preprocessing (cleaning and scaling)
- Step 3: Apply K-Means clustering to group customers
- Step 4: Train the model and save it as a .pkl file
- Step 5: Provide user input through the web interface
- Step 6: Predicts the nearest cluster.
- Step 7: Displays the predicted customer segment.

Based on prior learning, the system assigns customers to a specific group. Prediction occurs once training finishes completely. A predefined category emerges after analysis runs. The algorithm identifies patterns before sorting individuals. Following data review, classification takes place automatically. Outcome depends on earlier instructional phases.

3.5.3 Results Obtained

One way customers were sorted was by shared traits, forming clear groups. Through analysis, patterns emerged that defined separate segments. Some individuals fell into categories marked by similar behaviors. Groupings appeared when data points lined up in particular ways. Distinct profiles took shape once comparisons began. Certain trends set one collection apart from another. Not every person fit neatly, yet common features helped divide them. What stood out was how attributes aligned within each section

- High income – high spending customers
- Low income – low spending customers
- Moderate income – moderate spending customers

Without delay, outcomes appeared once the system classified each entry accurately. Patterns in how customers acted became clearer through divided groups.




Fig. 5. Customer Segmentation using K-Means (PCA Projection)




Fig. 6. Customer Segmentation System Dashboard Interface




Fig. 7. Predicted Customer Segment Output Interface

3.5.4 Validation of Results

Looking at how data grouped together helped confirm what the system found. These groupings made sense, matching how values spread across the set. Not every pattern stood out immediately, yet their structure followed expected lines.

For example:

Those earning more while also showing strong purchase patterns ended up in the same cluster

A portion of customers - those earning less while also spending little - grouped into their own segment. Not everyone fit neatly into broader patterns; some stood apart by financial behavior alone. Individuals living on tighter budgets often limited expenses further still. This distinct combination carved out a unique space within the data. Behavior like this emerged without prompting, simply through observation of transactions

Patterns in customer data are indeed recognized by the model, as shown here.

3.5.5 Verification of System

Through the web interface, various inputs were tested to confirm the system's behavior. When given comparable entries, it produced unchanged outputs. This pattern suggests the model performs steadily. Stability in results supports its dependability.

Quick responses marked the system's behavior, showing it worked well under test conditions. Throughout evaluations, operation remained stable without significant issues appearing.

Predict Customer
Enter customer attributes to classify into a segment

Customer Attributes

AGE	ANNUAL INCOME (\$)
25	45000
TOTAL SPENDING (\$)	RECENCY (DAYS)
2000	27
WEB PURCHASES	STORE PURCHASES
7	2
WEB VISITS / MONTH	
12	

Predict Segment →

Fig. 8. Customer Input Interface for Segmentation Prediction

3.5.6 Conclusion of Results

Despite its simplicity, the approach handles customer grouping well in practice. When combined with an online platform, K-Means proves useful for studying user patterns without complexity. By linking clustering techniques to accessible tools, analysis becomes more manageable over time.

3.6 Analysis

Looking at how well the method works reveals insights about its speed, output quality, and impact on sorting clients. From test outcomes, one sees clear patterns emerge when applying K-Means Clustering to separate buyer groups.

Among those clustered, individuals sharing traits appear close in grouping. Take, for instance, high earners who also spend more - they form a distinct segment. On the flip side, lower-income, low-spending customers land elsewhere. Pattern recognition stands out

here, showing how ties in data emerge naturally through this method.

Starting off differently each time helps clarity emerge when handling raw information before grouping it. Features gain balanced influence through adjustments like rescaling or removing errors, which stops stronger values from dominating. This balance means outcomes reflect actual patterns, not distortions caused by uneven ranges. Without consistent preparation steps, clusters might misrepresent how data truly connects.

Speed and efficiency mark strong points for the system. After training completes, outputs appear fast via the web platform. Such responsiveness fits scenarios needing rapid choices, including live or almost instant uses.

What stands out next is how straightforward and user-friendly the setup really is. Because the machine learning component connects directly to a browser platform, people without coding skills can still navigate it without trouble. That kind of accessibility turns it into something companies can actually use day to day.

Still, the study highlights certain constraints. Dependence on choosing the right cluster count (K) affects how well the model performs. Choosing a poor K value leads to weaker clustering outcomes. With only a narrow set of features and fixed data, the setup might miss key aspects of actual conditions.

Even with its drawbacks, the system works well enough. What stands out is how steadily the model groups data in sensible ways - showing clear patterns each time. Through this approach, machine learning proves useful when applied to actual challenges people face. Results stay reliable across different tests, which helps build confidence in the method.




Fig. 9. Customer Cluster Reference Description

3.7 Conclusion and Future work

Conclusion :

K-Means clustering is a machine learning technique that groups people based on things like their age, income, and spending. This helps businesses understand their customers better by finding patterns in data and making useful categories. This method turns raw data into useful information that can be used to make decisions. It is more accurate, faster, and automated than older methods. The system has a web interface that is easy to use, so even people who aren't tech-savvy can do segmentation quickly. Overall, it gives businesses accurate, reliable, and useful results that help them make better decisions based on data.

Future Work :

The system works very well, but there are still some gaps, and not all areas are guaranteed to get better. Future improvements might include adding real-time data so that the system can adjust to changes in customer behavior without having to be completely reprocessed. Other features, like location and customer preferences, could also be added to give more detailed information. If you have access to high-quality data, you can use advanced machine learning algorithms to make things even more accurate. Dynamic determination of the number of clusters during analysis, along with interactive visual interfaces, can make it easier to understand and more enjoyable for users. Also, connecting to business intelligence platforms and putting the software in the cloud can make it easier to use, more scalable, and more efficient. In general, these changes can make things more adaptable, stable, and useful in real-world business settings.

References :

- [1] R. Nugroho and M. Santoso, “*Customer Segmentation in Sales Transaction Data Using K-Means Clustering Algorithm*,” Journal of Intelligent Decision Support Systems, 2024.
- [2] Y. Syahra, A. Fadlil, and H. Yuliansyah, “*Customer segmentation using RFM and K-means clustering to support CRM in retail industry*,” SinkrOn: Jurnal dan Penelitian Teknik Informatika, vol. 9, no. 3, pp. 1120–1131, 2025.
- [3] M. Rumi, “*Customer Segmentation Using the K-Means Algorithm for Marketing Strategy Design*,” 2025.
- [4] X. Xu and Y. Tian, “*A Comprehensive Survey of Clustering Algorithms*,” Annals of Data Science, 2015.
- [5] R. Sharma and A. Gupta, “*Customer Segmentation Using K-Means Clustering Algorithm*,” International Journal of Creative Research Thoughts (IJCRT), 2024.
- [6] Dataset : <https://www.kaggle.com/datasets/vishakhdapat/customer-segmentation-clustering>
- [7] Some figures and diagrams generated using AI-assisted tools.